

# Bayesian Learning - Lab3

Jaskirat S Marar & Dinuke Jayaweera

5/3/2022

## GIBBS SAMPLER

```
# QUES 1: GIBBS SAMPLER
```

```
prec_data <- readRDS("precipitation.rds")
```


```
log_data <- log(prec_data)
```

```
par(mfrow = c(1,2))
```


```
hist(log_data, breaks = 40, prob = TRUE)
```

```
hist(prec_data, breaks = 40, prob = TRUE)
```

**Histogram of log\_data**



**Histogram of prec\_data**



## PART A

We first set up the priors and posteriors as follows:

PRIORS:

$$\mu \sim N(\mu_0, \tau_0^2)$$
$$\sigma^2 \sim \text{Inv} - \chi^2(\nu_0, \sigma_0^2)$$

FULL CONDITIONAL POSTERIOR:

$\mu$

$$\mu | \sigma^2, x \sim N(\mu_n, \tau_n^2)$$

where,  $\mu_n = w * (\bar{x}) + (1 - w) * \mu_0$ ,

$$\frac{1}{\tau_n^2} = \frac{1}{(\frac{n}{\sigma^2} + \frac{1}{\tau_0^2})},$$
$$w = \frac{\frac{n}{\sigma^2}}{(\frac{n}{\sigma^2} + \frac{1}{\tau_0^2})}$$

$\sigma^2$

$$\sigma^2 | \mu, x \sim \text{Inv} - \chi^2\left(\nu_n, \frac{\nu_0 \sigma_0^2 + \sum_{i=1}^n (x_i - \mu)^2}{n + \nu_0}\right)$$

where,  $\nu_n = \nu_0 + n$

We'll have to initialize some values for the parameters of the priors. For  $\mu_0$  &  $\sigma_0^2$  we choose the sample mean and sample variance. For  $\nu_0$  &  $\tau_0^2$  we choose to initialize them to 1 since the data is log normally distributed.

*#PART A*

*# setting up the initial values of the priors*

```
n = length(log_data)
```

*# for mu*

```
mu0 = mean(log_data)
```

```
tau20 = 1
```

*# for sigma*

```
sig20 = var(log_data)
```

```
nu0 = 1
```

```
set.seed(1234)
```

*#Gibbs Sampling function*

```
Gibbs_sampler <- function(steps = 500,  
                           mu_0 = mu0,  
                           tau2_0 = tau20,  
                           sig2_0 = sig20,  
                           nu_0 = nu0,  
                           n = length(log_data)){
```

*#initial value for gibbs sampler*

```
mu_i = rnorm(1, mean = mu_0, sd = sqrt(tau2_0))
```

```
sigma2_i = rinvchisq(1, df = nu_0, scale = sig2_0)
```

*#initialize result DF*

```
sample_results = data.frame("mu" = mu_i,
```

```

                                "sigma2" = sigma2_i)

#initialize loop over specified gibbs steps
for (i in 1:steps) {

  #gibbs step for mu
  w = n/sigma2_i / (n/sigma2_i + 1/tau2_0)
  mu_n = w * mean(log_data) + (1-w) * mu_0
  tau2_n = (n/sigma2_i + 1/tau2_0)^-1
  mu_i = rnorm(1, mean = mu_n, sd = sqrt(tau2_n))

  #gibbs step for sigma
  nu_n = nu_0 + n
  scale_term = (nu_0 * sigma2_i + sum((log_data - mu_i)^2))/nu_n
  sigma_i = rinvchisq(1, df = nu_n, scale = scale_term)

  #IF
  #IF_mu = 1 + 2*(sum(cor(sample_results[i,1], mu_i)))

  #save step values to DF
  sample_results = rbind.data.frame(sample_results, c(mu_i, sigma_i))

}

return(sample_results)


}

joint_posterior <- Gibbs_sampler(steps = 500)


#autocorrelation
par(mfrow = c(2,1))
mu_Gibbs <- acf(joint_posterior[,1])
sig2_Gibbs <- acf(joint_posterior[,2])

```

### Series joint\_posterior[, 1]



### Series joint\_posterior[, 2]



```
#inefficiency factor
IF_Gibbs_mu <- 1+2*sum(mu_Gibbs$acf[-1])
IF_Gibbs_sig2 <- 1+2*sum(sig2_Gibbs$acf[-1])


cat("\n", "IF of mu:", round(IF_Gibbs_mu,3))

##
## IF of mu: 1.147

cat("\n", "IF of sigma2:", round(IF_Gibbs_sig2,3))

##
## IF of sigma2: 0.421

#trace plot
par(mfrow = c(2,1))
plot(1:length(joint_posterior$mu), joint_posterior[,1], type = "l", col="blue", xlab = "steps", ylab = "mu")
plot(1:length(joint_posterior$sigma2), joint_posterior[,2], type = "l", col="blue", xlab = "steps", ylab = "sigma2")
```



From the trace plots it's very clear that the burn-in period is very short and that both chains converge very quickly.

## PART B

Since the data is log normally distributed and we've modelled the posterior on the  $\log(\text{data})$ , we can use the posterior draws and exponentiate them to compare to the original data.


```
# PART B

set.seed(1234)
post_pred <- c()
post_burnremoved <- joint_posterior[51:length(joint_posterior$mu),]
for (i in 1:(length(post_burnremoved$mu))){
  post_pred[i] = rnorm(1, mean = post_burnremoved[i,1], sd = sqrt(post_burnremoved[i,2]))
}

e_post_draws <- exp(post_pred)

plot(density(prec_data),
     col = "blue",
     ylim = c(0,0.2),
     xlim = c(-10,100),
     main = "Posterior Prediction vs Y")
lines(density(e_post_draws),
     col = "red")
```

## Posterior Prediction vs Y



# METROPOLIS RANDOM WALK

## PART A

```
# QUES 2: METROPOLIS RW

rm(list = ls())

data <- read.table("eBayNumberOfBidderData.dat", header = T)

# PART A

#drop intercept
data_nointercept <- data[,-2]

glm_model <- glm(nBids ~ ., family = poisson, data = data_nointercept)
summary(glm_model)

##
## Call:
## glm(formula = nBids ~ ., family = poisson, data = data_nointercept)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.5800  -0.7222  -0.0441   0.5269   2.4605
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  1.07244    0.03077  34.848 < 2e-16 ***
## PowerSeller -0.02054    0.03678  -0.558  0.5765
## VerifyID    -0.39452    0.09243  -4.268 1.97e-05 ***
## Sealed      0.44384    0.05056   8.778 < 2e-16 ***
## Minblem    -0.05220    0.06020  -0.867  0.3859
## MajBlem    -0.22087    0.09144  -2.416  0.0157 *
## LargNeg     0.07067    0.05633   1.255  0.2096
## LogBook    -0.12068    0.02896  -4.166 3.09e-05 ***
## MinBidShare -1.89410    0.07124 -26.588 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 2151.28  on 999  degrees of freedom
## Residual deviance:  867.47  on 991  degrees of freedom
## AIC: 3610.3
##
## Number of Fisher Scoring iterations: 5
```

Looking at the glm model output, we can see that the most significant features can be decided based on high  $\text{abs}(\beta)$  value with very small p-values. They are listed below:

1. MinBidShare
2. Intercept
3. Sealed

4. VerifyID
5. LogBook

## PART B

```
# PART B
#split response and features
response <- as.matrix(data$nBids)
covariates <- as.matrix(data[,2:10])
```

*Log-Likelihood:*

$$\text{Poisson PMF: } P(Y|X, \beta) = \frac{\lambda^Y}{Y!} \exp(-\lambda) \text{ where; } \lambda = \exp(X^T \beta)$$

$$P(y_i|x_i, \beta) = \prod_i \frac{\exp(y_i * x_i^T \beta) * \exp(-e^{x_i^T \beta})}{y_i!}$$

$$L(\beta|X, Y) = \prod_i \frac{\exp(y_i * x_i^T \beta) * \exp(-e^{x_i^T \beta})}{y_i!}$$

*taking logs both sides*

$$\log L(\beta|X, Y) = \sum_i \left( y_i * x_i^T \beta - e^{x_i^T \beta} - \log(y_i!) \right)$$

$$\log L(\beta|X, Y) = \sum_i \left( y_i * x_i^T \beta - e^{x_i^T \beta} \right)$$

*Zellner's g-Prior:*

$$\beta \sim N(0, 100 * (X^T X)^{-1})$$

*#log posterior*

```
logPost <- function(beta,
                     X = covariates,
                     Y = response){

  loglik <- sum(Y * (X %*% beta) - exp(X %*% beta))
  logprior <- dmvnorm(t(beta), mean = matrix(0, nrow = ncol(X)), sigma = (100*(solve(t(X) %*% X))), log
  return(loglik + logprior)

}
```

*#optimize log posterior*

```
opt_res <- optim(par = matrix(1, nrow = 9),
               fn = logPost,
               method = "BFGS",
               control = list(fnscale = -1),
               hessian = TRUE)

beta_mode <- opt_res$par
inv_jacobian <- -solve(opt_res$hessian)

rownames(beta_mode) <- colnames(covariates)
```



```

t(beta_mode)

##          Const PowerSeller  VerifyID    Sealed      Minblem      MajBlem      LargNeg
## [1,] 1.06981 -0.02050925 -0.3928761 0.4434933 -0.05235654 -0.2212323 0.07047041
##          LogBook MinBidShare
## [1,] -0.1202723 -1.892074

glm_model$coefficients

## (Intercept) PowerSeller  VerifyID      Sealed      Minblem      MajBlem
## 1.07244206 -0.02054076 -0.39451647 0.44384257 -0.05219829 -0.22087119
##      LargNeg      LogBook MinBidShare
## 0.07067246 -0.12067761 -1.89409664

```

We can see that the approximate posterior mode is very close to the MLE from the glm model.

## PART C

The metropolis algorithm will take any arbitrary posterior, which in this case that will be the log posterior & the tuning parameter  $c$ . The proposal density is given to be multivariate normal:

$$\theta_p | \theta^{i-1} \sim N(\theta^{i-1}, c.\Sigma)$$

*Parameter selection:*

*Target Density Prior:* For a fair comparison with part B of this question, we choose the same values as the prior parameters for beta as in the previous part.

*Proposal Density:* “ $\Sigma$  &  $\theta^{i-1}$  are initialized with the values calculated at the mode of the posterior in the previous step.

```

# PART C

# Metropolis Random Walk function
RMW_func <- function(target_density,
                      c,
                      theta_i_1,
                      sigma_proposal,
                      steps,
                      X,
                      Y){

  set.seed(12345)

  # initialize result matrix
  result <- matrix(t(theta_i_1), ncol = 9)

  accept = 0

  for (i in 1:steps) {

    #sample from the proposal distribution
    theta_p <- rmvnorm(1,
                      mean = as.vector(theta_i_1),
                      sigma = c*sigma_proposal)
  }
}

```

```

#calculate the ratio for the acceptance probability
ratio = ((target_density(as.vector(theta_p), X, Y))
          - (target_density(as.vector(theta_i_1), X, Y)))

#since target & proposal is defined in log form, we exponentiate to revert
ratio = exp(ratio)

#calculate alpha
alpha = min(1,ratio)

#draw from uniform
u <- runif(1)

#run test
if(u < alpha){
  accept = accept + 1
  theta_i_1 = theta_p
  result <- rbind(result, theta_i_1)
}
else {
  result = rbind(result, as.vector(theta_i_1))
}

}

return(list(result = result, acceptance = accept/steps))
}

# choose same params for the prior of target density
# choose data observed at mode for theta(i-1) & sigma for proposal density
sigma_proposal <- inv_jacobian
#theta_init <- beta_mode
theta_init <- matrix(rep(0.5,9), nrow = 9)

test_mrw <- RMW_func(target_density = logPost,
                     c = 0.6,
                     theta_i_1 = theta_init,
                     sigma_proposal = sigma_proposal,
                     steps = 5000,
                     X = covariates,
                     Y = response)

MRW_coeff <- test_mrw$result
MRW_coeff_means <- apply(MRW_coeff, 2, mean)
names(MRW_coeff_means) <- rownames(beta_mode)
MRW_coeff_means

##          Const  PowerSeller    VerifyID      Sealed      Minblem      MajBlem
## 1.053535037 -0.006066631 -0.384556715  0.446040824 -0.040732819 -0.211531073
##      LargNeg      LogBook  MinBidShare
## 0.081685121 -0.115261557 -1.855180954

```

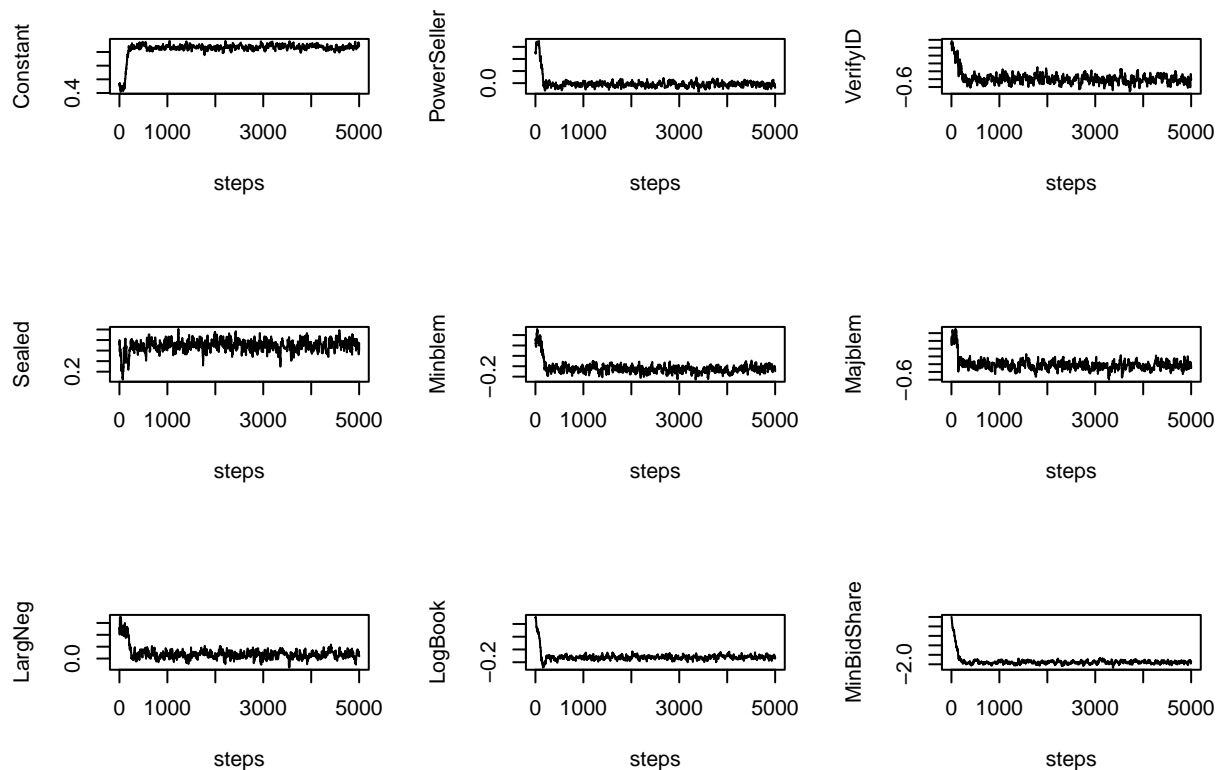
The choice of initial values of the  $\theta_{i-1}$  has an impact on how many steps we need for convergence in the random walk. When we chose the posterior mode from the previous analysis, then understandably we needed

very few steps for convergence, since the random walk algorithm doesn't have to move too much farther from the initial values to get to the actual mode. But if we initiate with very randomly chosen values for eg. 0, then we have to adjust accordingly for the burn in period. So, in order to show convergence of parameters properly, we chose a starting point of 0.5 for all coefficients. We are able to see convergence pretty soon on all chains, but to get a fairly stable result we need to run the sampler for 5000 steps, we are getting coefficient values in the neighborhood of what we got earlier from MLE and Bayesian analysis. We played around with different values of the tuning parameter 'c' to get a fair acceptance rate. With a value of 0.6, we are able to hit an acceptance rate of ~28% which is within the prescribed range. We plot the convergence of the coefficients as follows:

```
colnames(MRW_coeff) <- rownames(beta_mode)

par(mfrow = c(3,3))

plot(MRW_coeff[,1], type = 'l', ylab = "Constant", xlab = "steps")
plot(MRW_coeff[,2], type = 'l', ylab = "PowerSeller", xlab = "steps")
plot(MRW_coeff[,3], type = 'l', ylab = "VerifyID", xlab = "steps")
plot(MRW_coeff[,4], type = 'l', ylab = "Sealed", xlab = "steps")
plot(MRW_coeff[,5], type = 'l', ylab = "Minblem", xlab = "steps")
plot(MRW_coeff[,6], type = 'l', ylab = "Majblem", xlab = "steps")
plot(MRW_coeff[,7], type = 'l', ylab = "LargNeg", xlab = "steps")
plot(MRW_coeff[,8], type = 'l', ylab = "LogBook", xlab = "steps")
plot(MRW_coeff[,9], type = 'l', ylab = "MinBidShare", xlab = "steps")
```



## ## PART D

We'll now use the MCMC draws of the coefficients from the previous step as the  $\beta$  parameters to the given poisson regression model:

$$y_i|\beta \sim \text{Poisson}[\exp(x_i^T \beta)]$$

To find the distribution of  $\lambda$  as a linear combination, we will use the given features as  $X$  in the above model and the MCMC draws as  $\beta$ s. We'll then use the fitted  $\lambda$ s to draw from the poisson distribution to get a posterior prediction for the model.

**# PART D**

```
sample <- as.matrix(c(1, 1, 0, 1, 0, 1, 0, 1.2, 0.8), nrow = 9)
names(sample) <- rownames(beta_mode)
posterior_samples <- MRW_coeff[301:length(MRW_coeff[,1]),]


linear_predictions <- posterior_samples %*% sample
lambda_predictons <- exp(linear_predictions)

sample_predictions <- c()

set.seed(12345)
for (i in 1:length(lambda_predictons)){
  sample_predictions[i] <- rpois(1, lambda_predictons[i])
}

hist(sample_predictions, breaks = 20)
```

**Histogram of sample\_predictions**



```
cat("\n", "Probability of no bids on the test data:", length(which(sample_predictions == 0))/length(sample_predictions))
```

##

## Probability of no bids on the test data: 0.5052117

Based on our draws we can say that the probability that there are no bids for a listing with these features is about 50%.

# TIME SERIES MODELS

*AR(1) Process:*

$$x_t = \mu + \phi(x_{t-1} - \mu) + \varepsilon_t ; \text{ where } \varepsilon_t \widetilde{iid} N(0, \sigma^2)$$

## PART A

```
# QUES 3: TIME SERIES

rm(list = ls())

AR <- function(mu = 13,
               phi = c(-0.95, -0.5, 0, 0.25, 0.75, 0.95),
               sigma2 = 3,
               T = 300) {

  chains = data.frame(0, ncol = length(phi))

  for (i in 1:length(phi)){

    x_t <- mu
    chains[1,i] = x_t

    for (j in 2:T){

      x_t = mu + phi[i] * (x_t - mu) + rnorm(1, 0 , sqrt(sigma2))
      chains[j,i] = x_t

    }

  }


  colnames(chains) <- paste("phi_",phi)

  return(chains)

}

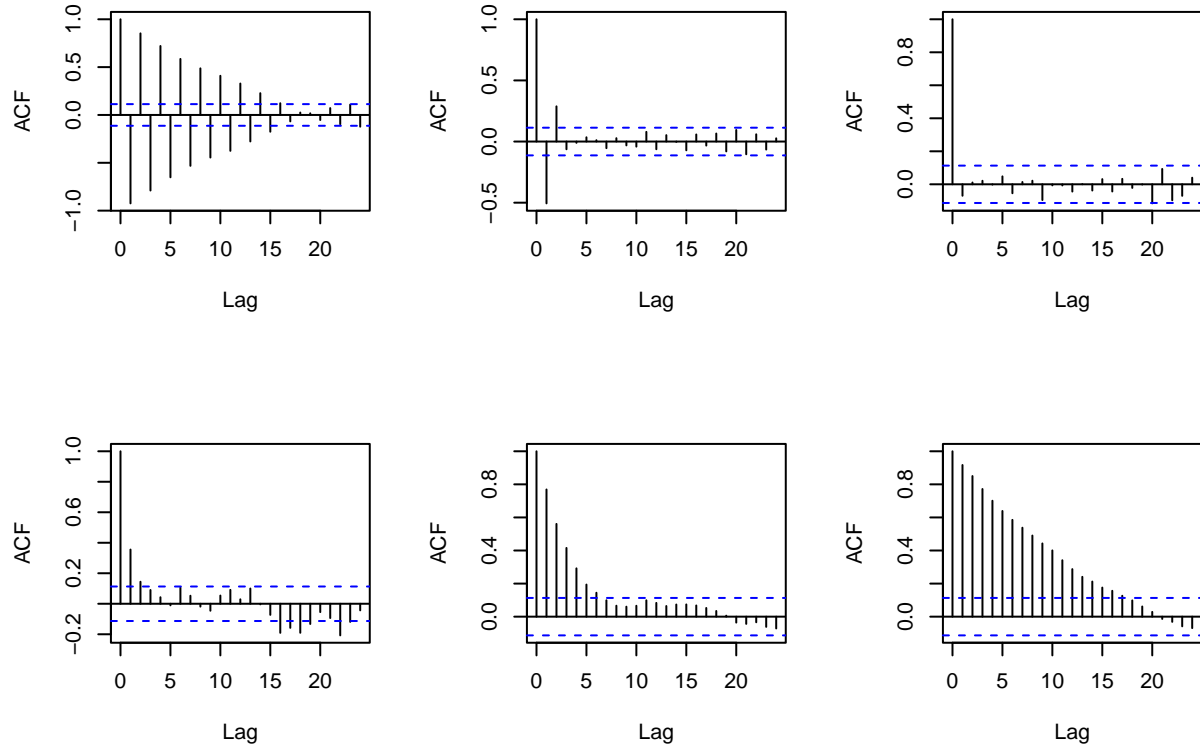
test_AR1 <- AR()

par(mfrow = c(3,3))
plot(y = test_AR1[,1], x = c(1:300), type = 'l', ylab = "phi = -0.95", xlab = "T")
plot(y = test_AR1[,2], x = c(1:300), type = 'l', ylab = "phi = -0.5", xlab = "T")
plot(y = test_AR1[,3], x = c(1:300), type = 'l', ylab = "phi = 0", xlab = "T")
plot(y = test_AR1[,4], x = c(1:300), type = 'l', ylab = "phi = 0.25", xlab = "T")
plot(y = test_AR1[,5], x = c(1:300), type = 'l', ylab = "phi = 0.75", xlab = "T")
plot(y = test_AR1[,6], x = c(1:300), type = 'l', ylab = "phi = 0.95", xlab = "T")
```



For lower values of  $\phi$  we can see that the error term of the AR(1) process i.e.  $\varepsilon_t$  has a much larger impact on the next prediction as the  $\mu$  dependent terms cancel out each others effect on the prediction, thus resulting in a negative correlation. But as we go towards a larger value of  $\phi$ , the first 2 terms dependent on  $\mu$  become additive and we see a stronger positive correlation in the chain for every subsequent prediction. To better depict whats going on with the chains, we plot the ACFs for them:

```
par(mfrow = c(2,3))
sapply(test_AR1, function(x) acf(x, main=""))
```



```
##      phi_ -0.95    phi_ -0.5    phi_ 0    phi_ 0.25    phi_ 0.75
## acf    numeric,25    numeric,25    numeric,25    numeric,25    numeric,25
## type   "correlation" "correlation" "correlation" "correlation" "correlation"
## n.used 300          300          300          300          300
## lag    numeric,25    numeric,25    numeric,25    numeric,25    numeric,25
## series "x"           "x"           "x"           "x"           "x"
## snames NULL         NULL         NULL         NULL         NULL
##      phi_ 0.95
## acf    numeric,25
## type   "correlation"
## n.used 300
## lag    numeric,25
## series "x"
## snames NULL
```

As we can infer from the ACF plots, for a positive value of the  $\phi$  with increasing lag, the autocorrelation keeps decreasing. Whereas for a negative value of  $\phi$  the autocorrelation keeps oscillating between consecutive lags. For  $\phi = 0$ , the successive values are independent of each other as the only variation is due to the normally distributed white noise.

## PART B

Now we are assuming the three parameters  $\mu$ ,  $\sigma^2$  and  $\phi$  to be unknown and are going to assume non-informative priors for them as follows:

1.  $\mu \sim N(0, 50)$  ; going with what we have used earlier for modelling  $\mu$  for a normal model
2.  $\sigma^2 \sim Inv - \chi^2(1, 10)$  ; since we used the inverse- $\chi^2$  distribution to model  $\sigma^2$  for a normal model



3.  $\phi \sim \text{Unif}(-1, 1)$  ; since we know the interval for phi

Since, it can be shown that if the white noise  $\varepsilon_t$  is gaussian then the AR(1) process for  $x_t$  is also gaussian with the following parameters:

$$x_t | x_{t-1} \sim N(\mu + \phi(x_{t-1} - \mu), \sigma_\varepsilon^2)$$

```
# PART B

model_AR <- AR(phi = c(0.2, 0.95))

StanModel = '
data {
  int<lower=0> T; // Number of observations
  vector[T] x;  // indicating the chain x over T obs
}
parameters {
  real mu;
  real<lower = 0> sigma2;
  real<lower = -1, upper = 1> phi;
}
model {
  mu ~ normal(0,50); // Normal with mean 0, st.dev. 50
  sigma2 ~ scaled_inv_chi_square(1,10); // Scaled-inv-chi2 with nu 1,sigma 10
  phi ~ uniform(-1,1);

  for(i in 2:T){
    x[i] ~ normal(mu + phi * (x[i-1] - mu), sqrt(sigma2));
  }
}'

#fit model for phi 0.2
fit_0.2 = stan(model_code = StanModel,
               data = list(x = model_AR$`phi_ 0.2`, T = 300),
               warmup = 1000,
               iter = 2000,
               chains = 4)

##
## SAMPLING FOR MODEL 'c4b76683c55ce82da98055a195c28b7c' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 0 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 1: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 1: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 1: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 1: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 1: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 1: Iteration:  1001 / 2000 [ 50%] (Sampling)
## Chain 1: Iteration:  1200 / 2000 [ 60%] (Sampling)
## Chain 1: Iteration:  1400 / 2000 [ 70%] (Sampling)
```

```

## Chain 1: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 1: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.15 seconds (Warm-up)
## Chain 1: 0.157 seconds (Sampling)
## Chain 1: 0.307 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'c4b76683c55ce82da98055a195c28b7c' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 0 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 2: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 2: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 2: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 2: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 2: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 2: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 2: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 2: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 2: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 2: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.132 seconds (Warm-up)
## Chain 2: 0.161 seconds (Sampling)
## Chain 2: 0.293 seconds (Total)
## Chain 2:
##
## SAMPLING FOR MODEL 'c4b76683c55ce82da98055a195c28b7c' NOW (CHAIN 3).
## Chain 3:
## Chain 3: Gradient evaluation took 0 seconds
## Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Chain 3: Adjust your expectations accordingly!
## Chain 3:
## Chain 3:
## Chain 3: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 3: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 3: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 3: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 3: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 3: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 3: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 3: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 3: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 3: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 3: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 3:

```

```

## Chain 3: Elapsed Time: 0.15 seconds (Warm-up)
## Chain 3:           0.156 seconds (Sampling)
## Chain 3:           0.306 seconds (Total)
## Chain 3:
##
## SAMPLING FOR MODEL 'c4b76683c55ce82da98055a195c28b7c' NOW (CHAIN 4).
## Chain 4:
## Chain 4: Gradient evaluation took 0 seconds
## Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Chain 4: Adjust your expectations accordingly!
## Chain 4:
## Chain 4:
## Chain 4: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 4: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 4: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 4: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 4: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 4: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 4: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 4: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 4: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 4: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 4: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 4:
## Chain 4: Elapsed Time: 0.162 seconds (Warm-up)
## Chain 4:           0.141 seconds (Sampling)
## Chain 4:           0.303 seconds (Total)
## Chain 4:

```

```

#fit model for phi 0.95
fit_0.95 = stan(model_code = StanModel,
                data = list(x = model_AR$`phi_ 0.95`, T = 300),
                warmup = 1000,
                iter = 2000,
                chains = 4)

```

```

##
## SAMPLING FOR MODEL 'c4b76683c55ce82da98055a195c28b7c' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 0 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 1: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 1: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 1: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 1: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 1: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 1: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 1: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 1: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 1: Iteration: 1600 / 2000 [ 80%] (Sampling)

```

```

## Chain 1: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.218 seconds (Warm-up)
## Chain 1: 0.163 seconds (Sampling)
## Chain 1: 0.381 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'c4b76683c55ce82da98055a195c28b7c' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 0 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 2: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 2: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 2: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 2: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 2: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 2: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 2: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 2: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 2: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 2: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.21 seconds (Warm-up)
## Chain 2: 0.142 seconds (Sampling)
## Chain 2: 0.352 seconds (Total)
## Chain 2:
##
## SAMPLING FOR MODEL 'c4b76683c55ce82da98055a195c28b7c' NOW (CHAIN 3).
## Chain 3:
## Chain 3: Gradient evaluation took 0 seconds
## Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Chain 3: Adjust your expectations accordingly!
## Chain 3:
## Chain 3:
## Chain 3: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 3: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 3: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 3: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 3: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 3: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 3: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 3: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 3: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 3: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 3: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 3:
## Chain 3: Elapsed Time: 0.242 seconds (Warm-up)

```

```

## Chain 3:          0.222 seconds (Sampling)
## Chain 3:          0.464 seconds (Total)
## Chain 3:
##
## SAMPLING FOR MODEL 'c4b76683c55ce82da98055a195c28b7c' NOW (CHAIN 4).
## Chain 4:
## Chain 4: Gradient evaluation took 0 seconds
## Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0 seconds.
## Chain 4: Adjust your expectations accordingly!
## Chain 4:
## Chain 4:
## Chain 4: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 4: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 4: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 4: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 4: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 4: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 4: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 4: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 4: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 4: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 4: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 4:
## Chain 4: Elapsed Time: 0.278 seconds (Warm-up)
## Chain 4:          0.269 seconds (Sampling)
## Chain 4:          0.547 seconds (Total)
## Chain 4:

## Warning: There were 1 divergent transitions after warmup. See
## https://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup
## to find out why this is a problem and how to eliminate them.

## Warning: Examine the pairs() plot to diagnose sampling problems

```

```

#PART B.i
#get posterior samples and means
post_samples_0.2 <- extract(fit_0.2)
post_mean_0.2 <- get_posterior_mean(fit_0.2)

post_samples_0.95 <- extract(fit_0.95)
post_mean_0.95 <- get_posterior_mean(fit_0.95)

cat("\n")

```

```

print("Posterior Means when phi = 0.2 :")

```

```

## [1] "Posterior Means when phi = 0.2 :"

```

```

cat("\n")

```

```

print(post_mean_0.2)

```

```

##          mean-chain:1 mean-chain:2 mean-chain:3 mean-chain:4 mean-all chains
## mu          12.9204957   12.9303981   12.9249563   12.9247568   12.9251517
## sigma2       3.3237206    3.3251022    3.3289650    3.3377813    3.3288922
## phi          0.3106549    0.3088885    0.3086545    0.3061329    0.3085827

```

```

## lp__    -330.7501262 -330.7647463 -330.7196701 -330.7139225    -330.7371163
cat("\n")

print("Posterior Means when phi = 0.95 :")

## [1] "Posterior Means when phi = 0.95 :"
cat("\n")

print(post_mean_0.95)

##          mean-chain:1 mean-chain:2 mean-chain:3 mean-chain:4 mean-all chains
## mu          11.7034128  12.5272756  12.2439620  12.1933107    12.166990
## sigma2       3.2775391   3.2854598   3.3100297   3.2613159    3.283586
## phi          0.9140362   0.9174577   0.9188501   0.9196601    0.917501
## lp__        -330.8182469 -330.9610729 -330.7634806 -331.0478086   -330.897652
post_samples_0.2_df <- data.frame(mu = post_samples_0.2$mu,
                                sigma2 = post_samples_0.2$sigma2,
                                phi = post_samples_0.2$phi)

CI_0.2 <- sapply(post_samples_0.2_df, function(x) quantile(x, probs=c(0.025, 0.975)))

post_samples_0.95_df <- data.frame(mu = post_samples_0.95$mu,
                                sigma2 = post_samples_0.95$sigma2,
                                phi = post_samples_0.95$phi)

CI_0.95 <- sapply(post_samples_0.95_df, function(x) quantile(x, probs=c(0.025, 0.975)))

cat("\n")

print("Posterior 95% CI when phi = 0.2 :")

## [1] "Posterior 95% CI when phi = 0.2 :"
cat("\n")

print(CI_0.2)

##          mu    sigma2    phi
## 2.5%  12.62712  2.826965 0.1921938
## 97.5% 13.22725  3.894716 0.4257022
cat("\n")

print("Posterior 95% CI when phi = 0.95 :")

## [1] "Posterior 95% CI when phi = 0.95 :"
cat("\n")

print(CI_0.95)

##          mu    sigma2    phi
## 2.5%   8.707761  2.803629 0.8600886
## 97.5% 16.173745  3.862484 0.9786845



```

Yes, we are able to estimate the true values of the parameters with reasonable width of CIs. But in case of the model for  $\phi = 0.95$ , the CI for  $\mu$  is very wide as compared to the other model.

*#PART B.ii.*







*#plotting the two models*

```
par(mfrow = c(2,1))
plot(y = model_AR[,1], x = c(1:300), type = 'l', ylab = "phi = 0.2", xlab = "T")
plot(y = model_AR[,2], x = c(1:300), type = 'l', ylab = "phi = 0.95", xlab = "T")
```




*#checking convergence of the samplers*

```
par(mfrow = c(3,2))
plot(post_samples_0.2$mu, type = 'l', ylab = "posterior Mu", main = "phi = 0.2")
plot(post_samples_0.95$mu, type = 'l', ylab = "posterior Mu", main = "phi = 0.95")
plot(post_samples_0.2$sigma2, type = 'l', ylab = "posterior Sigma2", main = "phi = 0.2")
plot(post_samples_0.95$sigma2, type = 'l', ylab = "posterior Sigma2", main = "phi = 0.95")
plot(post_samples_0.2$phi, type = 'l', ylab = "posterior phi", main = "phi = 0.2")
plot(post_samples_0.95$phi, type = 'l', ylab = "posterior phi", main = "phi = 0.95")
```



```
#joint posterior
par(mfrow = c(1,2))
plot(x = post_samples_0.2_df$mu,
     y = post_samples_0.2_df$phi,
     main = "phi = 0.2",
     xlab = "mu",
     ylab = "phi")
plot(x = post_samples_0.95_df$mu,
     y = post_samples_0.95_df$phi,
     main = "phi = 0.95",
     xlab = "mu",
     ylab = "phi")
```





All the parameter chains are converging pretty quickly. From the joint posterior it is evident that the estimate for  $\mu$  is more accurate for the model with  $\phi = 0.2$  than the other model. An explanation for this can be that since the sample mean in an AR(1) process with a shifted mean (as is the case here), is  $\sim N(0, \sigma^2/(1 - \phi)^2)$ , the CI for the sample mean is then given by:

$$C.I.,_{.95}(\mu) = \bar{x}_n \pm \frac{1.96}{\sqrt{n}} \left( \frac{\sigma}{|1 - \phi|} \right)$$

From this expression we can concur that the CI for the mean will be much wider when the  $\phi$  is small. From the joint posterior plots we can draw a similar inference that as the value of  $\phi$  starts approaching 1, the stationary properties start getting lost and the distribution of  $\mu$  becomes very wide.

## APPENDIX: CODE

```
knitr::opts_chunk$set(echo = TRUE)
library(mvtnorm)
library(LaplacesDemon)
library(rstan)

# QUES 1: GIBBS SAMPLER

prec_data <- readRDS("precipitation.rds")

log_data <- log(prec_data)

par(mfrow = c(1,2))
hist(log_data, breaks = 40, prob = TRUE)
hist(prec_data, breaks = 40, prob = TRUE)

#PART A

# setting up the initial values of the priors

n = length(log_data)

# for mu
mu0 = mean(log_data)
tau20 = 1
# for sigma
sig20 = var(log_data)
nu0 = 1

set.seed(1234)

#Gibbs Sampling function
Gibbs_sampler <- function(steps = 500,
                          mu_0 = mu0,
                          tau2_0 = tau20,
                          sig2_0 = sig20,
                          nu_0 = nu0,
                          n = length(log_data)){

  #initial value for gibbs sampler
  mu_i = rnorm(1, mean = mu_0, sd = sqrt(tau2_0))
  sigma2_i = rinvchisq(1, df = nu_0, scale = sig2_0)

  #initialize result DF
  sample_results = data.frame("mu" = mu_i,
                              "sigma2" = sigma2_i)

  #initialize loop over specified gibbs steps
  for (i in 1:steps) {

    #gibbs step for mu
    w = n/sigma2_i / (n/sigma2_i + 1/tau2_0)
```

```

mu_n = w * mean(log_data) + (1-w) * mu_0
tau2_n = (n/sigma2_i + 1/tau2_0)^-1
mu_i = rnorm(1, mean = mu_n, sd = sqrt(tau2_n))

#gibbs step for sigma
nu_n = nu_0 + n
scale_term = (nu_0 * sigma2_i + sum((log_data - mu_i)^2))/nu_n
sigma_i = rinvchisq(1, df = nu_n, scale = scale_term)

#IF
#IF_mu = 1 + 2*(sum(cor(sample_results[i,1], mu_i)))

#save step values to DF
sample_results = rbind.data.frame(sample_results, c(mu_i, sigma_i))

}

return(sample_results)

}

joint_posterior <- Gibbs_sampler(steps = 500)

#autocorrelation
par(mfrow = c(2,1))
mu_Gibbs <- acf(joint_posterior[,1])
sig2_Gibbs <- acf(joint_posterior[,2])

#inefficiency factor
IF_Gibbs_mu <- 1+2*sum(mu_Gibbs$acf[-1])
IF_Gibbs_sig2 <- 1+2*sum(sig2_Gibbs$acf[-1])

cat("\n", "IF of mu:", round(IF_Gibbs_mu,3))
cat("\n", "IF of sigma2:", round(IF_Gibbs_sig2,3))

#trace plot
par(mfrow = c(2,1))
plot(1:length(joint_posterior$mu), joint_posterior[,1], type = "l", col="blue", xlab = "steps", ylab = "mu")
plot(1:length(joint_posterior$sigma2), joint_posterior[,2], type = "l", col="blue", xlab = "steps", ylab = "sigma2")

# PART B

set.seed(1234)
post_pred <- c()
post_burnremoved <- joint_posterior[51:length(joint_posterior$mu),]
for (i in 1:(length(post_burnremoved$mu))){
  post_pred[i] = rnorm(1, mean = post_burnremoved[i,1], sd = sqrt(post_burnremoved[i,2]))
}

e_post_draws <- exp(post_pred)

plot(density(prec_data),

```

```

    col = "blue",
    ylim = c(0,0.2),
    xlim = c(-10,100),
    main = "Posterior Prediction vs Y")
lines(density(e_post_draws),
      col = "red")

# QUES 2: METROPOLIS RW

rm(list = ls())

data <- read.table("eBayNumberOfBidderData.dat", header = T)

# PART A

#drop intercept
data_nointercept <- data[,-2]

glm_model <- glm(nBids ~ ., family = poisson, data = data_nointercept)
summary(glm_model)

# PART B

#split response and features
response <- as.matrix(data$nBids)
covariates <- as.matrix(data[,2:10])

#log posterior

logPost <- function(beta,
                     X = covariates,
                     Y = response){

  loglik <- sum(Y * (X %*% beta) - exp(X %*% beta))
  logprior <- dmvnorm(t(beta), mean = matrix(0, nrow = ncol(X)), sigma = (100*(solve(t(X) %*% X))), log
  return(loglik + logprior)
}

#optimize log posterior

opt_res <- optim(par = matrix(1, nrow = 9),
                fn = logPost,
                method = "BFGS",
                control = list(fnscale = -1),
                hessian = TRUE)

beta_mode <- opt_res$par
inv_jacobian <- -solve(opt_res$hessian)

rownames(beta_mode) <- colnames(covariates)

t(beta_mode)

```

```

glm_model$coefficients
# PART C

# Metropolis Random Walk function
RMW_func <- function(target_density,
                      c,
                      theta_i_1,
                      sigma_proposal,
                      steps,
                      X,
                      Y){

  set.seed(12345)

  # initialize result matrix
  result <- matrix(t(theta_i_1), ncol = 9)

  accept = 0

  for (i in 1:steps) {

    #sample from the proposal distribution
    theta_p <- rmvnorm(1,
                      mean = as.vector(theta_i_1),
                      sigma = c*sigma_proposal)

    #calculate the ratio for the acceptance probability
    ratio = ((target_density(as.vector(theta_p), X, Y))
             - (target_density(as.vector(theta_i_1), X, Y)))

    #since target & proposal is defined in log form, we exponentiate to revert
    ratio = exp(ratio)

    #calculate alpha
    alpha = min(1,ratio)

    #draw from uniform
    u <- runif(1)

    #run test
    if(u < alpha){
      accept = accept + 1
      theta_i_1 = theta_p
      result <- rbind(result, theta_i_1)
    }
    else {
      result = rbind(result, as.vector(theta_i_1))
    }
  }

  return(list(result = result, acceptance = accept/steps))

```

```

}
# choose same params for the prior of target density
# choose data observed at mode for theta(i-1) & sigma for proposal density
sigma_proposal <- inv_jacobian
#theta_init <- beta_mode
theta_init <- matrix(rep(0.5,9), nrow = 9)

test_mrw <- RMW_func(target_density = logPost,
                     c = 0.6,
                     theta_i_1 = theta_init,
                     sigma_proposal = sigma_proposal,
                     steps = 5000,
                     X = covariates,
                     Y = response)

MRW_coeff <- test_mrw$result
MRW_coeff_means <- apply(MRW_coeff, 2, mean)
names(MRW_coeff_means) <- rownames(beta_mode)
MRW_coeff_means

colnames(MRW_coeff) <- rownames(beta_mode)

par(mfrow = c(3,3))

plot(MRW_coeff[,1], type = 'l', ylab = "Constant", xlab = "steps")
plot(MRW_coeff[,2], type = 'l', ylab = "PowerSeller", xlab = "steps")
plot(MRW_coeff[,3], type = 'l', ylab = "VerifyID", xlab = "steps")
plot(MRW_coeff[,4], type = 'l', ylab = "Sealed", xlab = "steps")
plot(MRW_coeff[,5], type = 'l', ylab = "Minblem", xlab = "steps")
plot(MRW_coeff[,6], type = 'l', ylab = "Majblem", xlab = "steps")
plot(MRW_coeff[,7], type = 'l', ylab = "LargNeg", xlab = "steps")
plot(MRW_coeff[,8], type = 'l', ylab = "LogBook", xlab = "steps")
plot(MRW_coeff[,9], type = 'l', ylab = "MinBidShare", xlab = "steps")

# PART D

sample <- as.matrix(c(1, 1, 0, 1, 0, 1, 0, 1.2, 0.8), nrow = 9)
names(sample) <- rownames(beta_mode)
posterior_samples <- MRW_coeff[301:length(MRW_coeff[,1]),]

linear_predictions <- posterior_samples %*% sample
lambda_predictons <- exp(linear_predictions)

sample_predictions <- c()

set.seed(12345)
for (i in 1:length(lambda_predictons)){
  sample_predictions[i] <- rpois(1, lambda_predictons[i])
}

hist(sample_predictions, breaks = 20)

cat("\n", "Probability of no bids on the test data:",length(which(sample_predictions == 0))/length(sample_predictions))

```

```

# QUES 3: TIME SERIES

rm(list = ls())

AR <- function(mu = 13,
               phi = c(-0.95, -0.5, 0, 0.25, 0.75, 0.95),
               sigma2 = 3,
               T = 300) {

  chains = data.frame(0, ncol = length(phi))

  for (i in 1:length(phi)){

    x_t <- mu
    chains[1,i] = x_t

    for (j in 2:T){

      x_t = mu + phi[i] * (x_t - mu) + rnorm(1, 0, sqrt(sigma2))
      chains[j,i] = x_t

    }

  }

  colnames(chains) <- paste("phi_",phi)

  return(chains)

}

test_AR1 <- AR()

par(mfrow = c(3,3))
plot(y = test_AR1[,1], x = c(1:300), type = 'l', ylab = "phi = -0.95", xlab = "T")
plot(y = test_AR1[,2], x = c(1:300), type = 'l', ylab = "phi = -0.5", xlab = "T")
plot(y = test_AR1[,3], x = c(1:300), type = 'l', ylab = "phi = 0", xlab = "T")
plot(y = test_AR1[,4], x = c(1:300), type = 'l', ylab = "phi = 0.25", xlab = "T")
plot(y = test_AR1[,5], x = c(1:300), type = 'l', ylab = "phi = 0.75", xlab = "T")
plot(y = test_AR1[,6], x = c(1:300), type = 'l', ylab = "phi = 0.95", xlab = "T")

par(mfrow = c(2,3))
sapply(test_AR1, function(x) acf(x, main=""))

# PART B

model_AR <- AR(phi = c(0.2, 0.95))

StanModel = '
data {
  int<lower=0> T; // Number of observations
  vector[T] x; // indicating the chain x over T obs

```

```

}
parameters {
  real mu;
  real<lower = 0> sigma2;
  real<lower = -1, upper = 1> phi;
}
model {
  mu ~ normal(0,50); // Normal with mean 0, st.dev. 50
  sigma2 ~ scaled_inv_chi_square(1,10); // Scaled-inv-chi2 with nu 1,sigma 10
  phi ~ uniform(-1,1);

  for(i in 2:T){
    x[i] ~ normal(mu + phi * (x[i-1] - mu), sqrt(sigma2));
  }
}'

#fit model for phi 0.2
fit_0.2 = stan(model_code = StanModel,
               data = list(x = model_AR$`phi_ 0.2`, T = 300),
               warmup = 1000,
               iter = 2000,
               chains = 4)

#fit model for phi 0.95
fit_0.95 = stan(model_code = StanModel,
                data = list(x = model_AR$`phi_ 0.95`, T = 300),
                warmup = 1000,
                iter = 2000,
                chains = 4)

#PART B.i
#get posterior samples and means
post_samples_0.2 <- extract(fit_0.2)
post_mean_0.2 <- get_posterior_mean(fit_0.2)

post_samples_0.95 <- extract(fit_0.95)
post_mean_0.95 <- get_posterior_mean(fit_0.95)

cat("\n")
print("Posterior Means when phi = 0.2 :")
cat("\n")
print(post_mean_0.2)

cat("\n")
print("Posterior Means when phi = 0.95 :")
cat("\n")
print(post_mean_0.95)

post_samples_0.2_df <- data.frame(mu = post_samples_0.2$mu,

```



```

        sigma2 = post_samples_0.2$sigma2,
        phi = post_samples_0.2$phi)

CI_0.2 <- sapply(post_samples_0.2_df, function(x) quantile(x, probs=c(0.025, 0.975)))

post_samples_0.95_df <- data.frame(mu = post_samples_0.95$mu,
        sigma2 = post_samples_0.95$sigma2,
        phi = post_samples_0.95$phi)

CI_0.95 <- sapply(post_samples_0.95_df, function(x) quantile(x, probs=c(0.025, 0.975)))

cat("\n")
print("Posterior 95% CI when phi = 0.2 :")
cat("\n")
print(CI_0.2)

cat("\n")
print("Posterior 95% CI when phi = 0.95 :")
cat("\n")
print(CI_0.95)

#PART B.ii.

#plotting the two models
par(mfrow = c(2,1))
plot(y = model_AR[,1], x = c(1:300), type = 'l', ylab = "phi = 0.2", xlab = "T")
plot(y = model_AR[,2], x = c(1:300), type = 'l', ylab = "phi = 0.95", xlab = "T")

#checking convergence of the samplers
par(mfrow = c(3,2))
plot(post_samples_0.2$mu, type = 'l', ylab = "posterior Mu", main = "phi = 0.2")
plot(post_samples_0.95$mu, type = 'l', ylab = "posterior Mu", main = "phi = 0.95")
plot(post_samples_0.2$sigma2, type = 'l', ylab = "posterior Sigma2", main = "phi = 0.2")
plot(post_samples_0.95$sigma2, type = 'l', ylab = "posterior Sigma2", main = "phi = 0.95")
plot(post_samples_0.2$phi, type = 'l', ylab = "posterior phi", main = "phi = 0.2")
plot(post_samples_0.95$phi, type = 'l', ylab = "posterior phi", main = "phi = 0.95")

#joint posterior
par(mfrow = c(1,2))
plot(x = post_samples_0.2_df$mu,
     y = post_samples_0.2_df$phi,
     main = "phi = 0.2",
     xlab = "mu",
     ylab = "phi")
plot(x = post_samples_0.95_df$mu,
     y = post_samples_0.95_df$phi,
     main = "phi = 0.95",
     xlab = "mu",
     ylab = "phi")

```